ContinuumBench: Independent Research Foundation

Samarth Annigeri

13 October 2026

Abstract

This document is a report skeleton, prepared before research observations exist. It records the intended argument and evidence requirements. It is not a completed study, an MSc thesis, or a substitute approved by an admissions committee.

1 Introduction

We investigate how organisations can make generative-AI deployment choices auditable. The project-specific research questions are recorded in the repository protocol.

2 Background and related work

The primary-source register distinguishes abstract-level reading from full-text review. The final report requires at least 25 verified references. That review is unfinished; no filler citations are included here.

3 Method

We separate design assumptions, synthetic software fixtures, measured observations and estimates. A protocol must be frozen before research runs. Human annotations must be authored by a human and retained separately from generated artefacts.

4 Implementation

The repository contains an initial executable vertical slice and quality-gate configuration. Passing software checks establishes behaviour within tested conditions, not stakeholder acceptance or model quality.

5 Results

Pending. There are no inference, elicitation or practitioner-study results. Any research table must be generated from recorded raw data.

6 Discussion

The design is intended to support later empirical work. Claims of benefit require comparison baselines, uncertainty estimates and evidence of generalisation.

7 Threats to validity

Construct validity: synthetic cases may miss real stakeholder concerns. Internal validity: software verification cannot establish model quality. External validity: one candidate's hardware and authored scenarios cannot represent European organisations. Conclusion validity: no empirical comparison can be made before data collection.

8 Future work

We will refine the case with practitioners, validate instruments, compare alternatives and test transfer across organisations. This is a proposed research agenda, not evidence of execution.

9 Reproducibility statement

Use the exact commit and lockfile. Research results are absent. The repository documents which remaining steps require annotation, hardware access or source verification.

10 AI-assistance statement

AI coding assistants produced this foundation from Samarth's brief. His annotations, benchmark runs and final review are pending. The project has no institutional endorsement.